

PRE-DEMO INVESTOR PITCH

CredRisk.AI

An underwriter workbench that helps SME lenders move from scattered review to explainable, human-controlled credit decisions.

- SME lending
- Risk scoring
- AI second review
- Audit trail

A-F

Clear risk grades mapped to model recommendations.

RF

Random Forest baseline with validation metrics.

LLM

Optional second reviewer, not an automatic final decision.

SME underwriting is too manual for the risk hiding inside the file.

Analysts spend time collecting context before they can judge risk.

Documents, financial ratios, fraud signals, portfolio exposure, and final decisions often live in separate tools.

Slow triage

High-risk cases wait in the same queue as clean applications.

Low transparency

Analysts need to explain why a score moved, not just see a number.

Weak governance

Model output, AI advice, and human final decisions must stay separate.

THE SOLUTION

One workbench for scoring, explaining, reviewing, and monitoring SME loan applications.

1

Score

RF model produces an application risk score and A-F grade.

2

Explain

Deterministic rationale and SHAP-style drivers make the score readable.

3

Review

The analyst records a final decision separately from the model.

4

Monitor

Queues, portfolio risk, and audit history keep operations visible.

CredRisk.AI gives lenders speed without losing control.

Operational speed

Assigned tasks, workboard filters, bulk rejection, and case handoff reduce review friction.

Explainable confidence

Risk drivers translate raw model signals into banker-readable rationale.

Human accountability

The tool recommends; Ms. Cooper decides, documents, and routes supervisor review when needed.

AI where useful

Hosted or local LLMs can challenge the baseline with a qualitative second review.

Bank-grade audit

Session decisions, memo generation, and audit history support governance.

Scalable foundation

The MVP can connect to PSD2, accounting, registry, KYB, and document pipelines.

A compact MVP that maps the business plan into working software.

1

Streamlit UI

Home, workspace, operations desk, dashboards, settings, and support.

2

Data pipeline

Synthetic SME portfolio, derived cash-flow, document, KYB, and anomaly signals.

3

RF model

Scikit-learn pipeline with preprocessing, Random Forest, A-F grading, and metrics.

4

Explainability

Deterministic rationale first, plus optional OpenAI or local LLM second review.

What to watch when the product demo starts.

- 1

Home

Ms. Cooper starts from tasks, Slack updates, and today's calendar.
- 2

Personal Workspace

Score an assigned SME case and read the deterministic decision rationale.
- 3

Case Review

Save a final analyst decision, separate from the RF recommendation.
- 4

Operations Desk

Filter the team workboard and apply bulk action where appropriate.
- 5

Risk Dashboard

Monitor exposure, manual reviews, compliance queue, and audit history.
- 6

AI Explainability

Run an optional LLM second review and compare it against the RF baseline.

The MVP is intentionally focused; the next layer is real data integration and governance hardening.

Scale prerequisites

Production data connectors, model monitoring, role-based access, document ingestion, secure secrets, and deployment controls.

- PSD2
- Accounting APIs
- KYB registry
- Document review

Managed risks

False positives, AI hallucination, data leakage, model drift, and over-reliance on automation are controlled through human review and auditability.

- Human final decision
- LLM fallback
- Validation metrics
- Audit trail

Built as a working MVP with a clear path to lender deployment.

Product works

The app implements the core value loop: scoring, explainability, analyst workflow, bulk operations, and monitoring.

Easy to run

The workspace includes setup docs, a demo guide, requirements, and a Windows launcher for fast handoff.

AI governed

Agent instructions and acceptance checks keep the build aligned with risk policy, page structure, and user workflow.

NOW THE DEMO

CredRisk.AI turns SME underwriting from a file chase into an explainable decision workflow.

The demo will show the product live: score a case, explain the score, record the human decision, review the portfolio, and run an AI second opinion.
